

# EECS 545: Machine Learning

## Lecture 1. Introduction

Honglak Lee & Michał Dereziński

1/5/2022



# Outline

- Administrative
- What is machine learning?

# Teaching staff

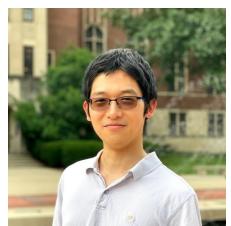
Instructors:

- Honglak Lee ([honglak@umich.edu](mailto:honglak@umich.edu), Office: CSE 3773)
- Michał Dereziński ([derezin@umich.edu](mailto:derezin@umich.edu), Office: CSE 2649)



Graduate Student Instructors:

- Aabhaas Vaish ([aabhaas@umich.edu](mailto:aabhaas@umich.edu))
- Junghwan Kim ([kimjhj@umich.edu](mailto:kimjhj@umich.edu))
- Kevin Wang ([musicer@umich.edu](mailto:musicer@umich.edu))
- Sudeep Katakol ([sudeepk@umich.edu](mailto:sudeepk@umich.edu))



Junghwan

Sudeep

Aabhaas

Kevin

**Please contact us via [Piazza](#) or [eeecs545.w22@gmail.com](mailto:eeecs545.w22@gmail.com)** (please avoid sending to personal emails unless it's really necessary)

- For office hours, please see future announcements and [calendar](#)

# About this course

- Graduate-level introduction of machine learning
- Foundations of machine learning
  - Mathematical derivation, Implementation of the algorithms, applications
- Topics
  - supervised learning
  - unsupervised learning
  - reinforcement learning
- Other topics
  - deep learning, learning theory, probabilistic models, sparsity and feature selection, Bayesian techniques, ensemble methods

# About this course

- Cover practical applications of machine learning
  - computer vision, data mining, speech recognition, text processing, robot perception and control, etc.
- Our goal is to help you to
  - Understand fundamentals of machine learning
  - Learn technical details of ML algorithms
  - Learn how to implement some important algorithms
  - Use machine learning algorithms for your research and applications.

# Textbooks

There will be no official textbook for the course. However, the following materials will be helpful (all of them are available online):

- Chris Bishop, “[Pattern Recognition and Machine Learning](#)”. Springer, 2007.
- Kevin Murphy, “Machine Learning: A Probabilistic Perspective”, 2012.
- David Barber, “[Bayesian Reasoning and Machine Learning](#)”, 2017.
- Ian Goodfellow and Yoshua Bengio and Aaron Courville, “[Deep Learning](#)”,
- Hastie, Tibshirani, Fiedman, “[Elements of Statistical Learning](#)”. Springer, 2010.
- Sutton and Barto, "[Reinforcement Learning: An Introduction](#)," MIT Press, 2018.
- (optional) Boyd and Vandenberghe, "[Convex Optimization](#)," Cambridge University Press, 2004.
- (optional) Mackay, “[Information Theory, Inference, and Learning Algorithms](#)”. Cambridge University Press. 2003.

# Prerequisites

- Undergrad linear algebra (e.g., MATH 217, MATH 417) OR graduate matrix methods courses (EECS 505 / EECS 551) which have several relevant linear algebra concepts
- Multivariate calculus
- Undergrad probability and statistics (e.g., EECS 301)
- Programming skills (equivalent to EECS 280, EECS 281, and experience in Python)
  - Nontrivial level of programming is required.
- NOTE: If you **have not** taken **at least two** of linear algebra, multivariate calculus, and probability courses, it is **strongly recommended** that you finish them first before taking this course.

# Grading policy

- Homework: 30%
- Midterm: 30% (tentative date: early April)
  - [Honor code](#)
- Project: 40%
  - progress report (10%)
  - final report (30%)
- Note: There will be no final exam.
- Extra credits: Up to 3% may be awarded for participation in class and piazza, as well as quizzes (details TBD).
  - Note: anonymous posts/comments are not counted

# Language of Choice: **Python**

- **Python** is a great language overall for machine learning, with modern libraries and excellent online tutorials for various tasks
- We will use **Python 3 (3.6+)** throughout the course.
  - We will be using popular libraries, such as **NumPy**, **Matplotlib**, and **PyTorch**.
- There will be a tutorial session (date: TBD)

# Homework

- There will be 6 problem sets.
  - **The best 5 out of 6 scores will be used for final grading.**
- Goal: strengthen the understanding of the fundamental concepts, mathematical formulations, algorithms, detailed implementations, and the applications.
- The problem sets will also include programming assignments to implement algorithms covered in the class.
- Homework #1 will be out by next Tuesday (1/11) – due 1/25, 11:55 pm via Gradescope (linked to Canvas).

# Late days

- 3 maximum late days per assignment (except homework #1)
  - No homework will be accepted 3 days after the due date.
  - Exception: due to Drop/Add deadline (Jan 25), homework #1 deadline will be Jan 25 with no late days (hard deadline; no late days will be allowed)
- Total 9 late days allowed.
- After using up all late days, your assignment will be penalized by 20% from your scores.

# Study group

- Form your study group early on!
  - Up to five people are allowed.
- For homework, you may discuss between the study group members, but you should write your own solution independently.
- In the homework submissions, you must put:
  - the names of other people you collaborated.
  - submission time.
- Please start on homework early. (Warning: cramming will not work!)

# Honor Code

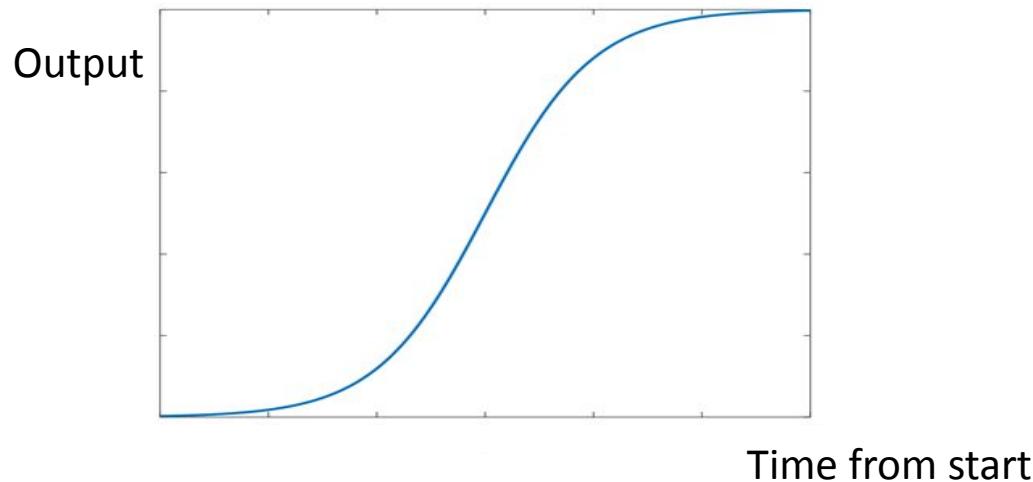
- No cheating of any kind shall be made.
- Previously, several incidents of violation of honor (plagiarism in homework submissions and cheatings in exams) were reported.
  - You will get into a **big trouble** if you violate the honor code. You may fail the course.
- Please review the UofM [Honor Code](#).

# Course Project

- Scope
  - develop new algorithms and theory in machine learning
  - apply existing algorithms to new problems
  - apply to your own problems of interest
- Milestones (tentative)
  - Project proposal due: February 1, 2022
  - Project milestone report due: March 8, 2022
  - Final project poster presentation: April 21, 2022
  - Final project report due: April 28, 2022 (No Late Days)
- Requirement
  - Write a 8-page paper
  - Submit the final code
  - Give a poster presentation
- Evaluation is based on:
  - novelty, technical quality, significance, and presentation quality of the project.
- More information available [here](#) (draft).

# Course Project

- 4 or 5 people can form a project group.
- Talk to instructor if you want to get suggestions about project topics.
- Start early! (form your group and start working)



# Expected Workload

- While EECS 545 is a 3 unit course, the expected workload is comparable to upper-level 4 unit courses.
- Previously, students who took EECS 545 reported the relative workload (Q891) as “lighter” (~15%), “typical” (~25%), “heavier” (~30%), or “much heavier” (~30%).
- You can expect to learn a lot of things but the workload could feel heavy or very heavy if you don’t have sufficient background in math (e.g., linear algebra, calculus, probabilities) and/or programming (e.g., nontrivial programming background).
- If you are concerned, please talk to us.

# Other Information

- Review sessions
  - Will hold review sessions on background materials (linear algebra, probability, Python.) Dates TBD
- Midterm
  - Tentatively scheduled for April 7, 2022
- Beginning-of-course survey:  
<https://forms.gle/rw2RCfY3GvCekZ646>

Any questions?

# Outline

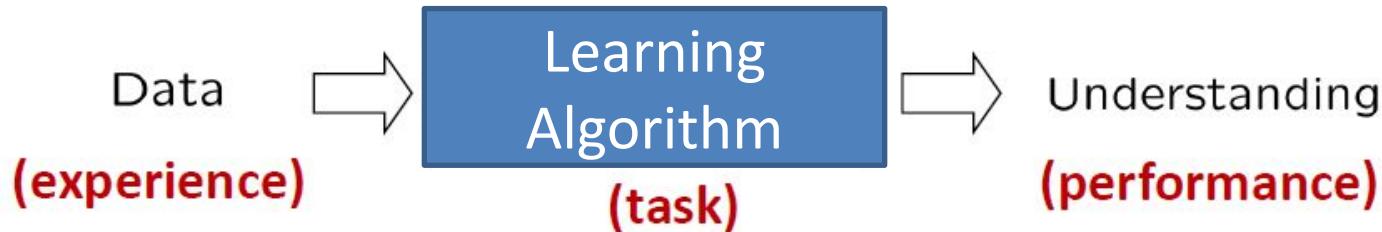
- Administrative
- What is machine learning?

# Definition of Machine Learning

- Formal definition (Mitchell 1997): A computer program **A** is said to **learn from experience E** with respect to some class of tasks T and performance measure P, if its performance at tasks in **T**, as measured by **P**, improves with experience **E**.

# Informal definition

- Algorithms that improve their prediction performance at some task with experience (or data).



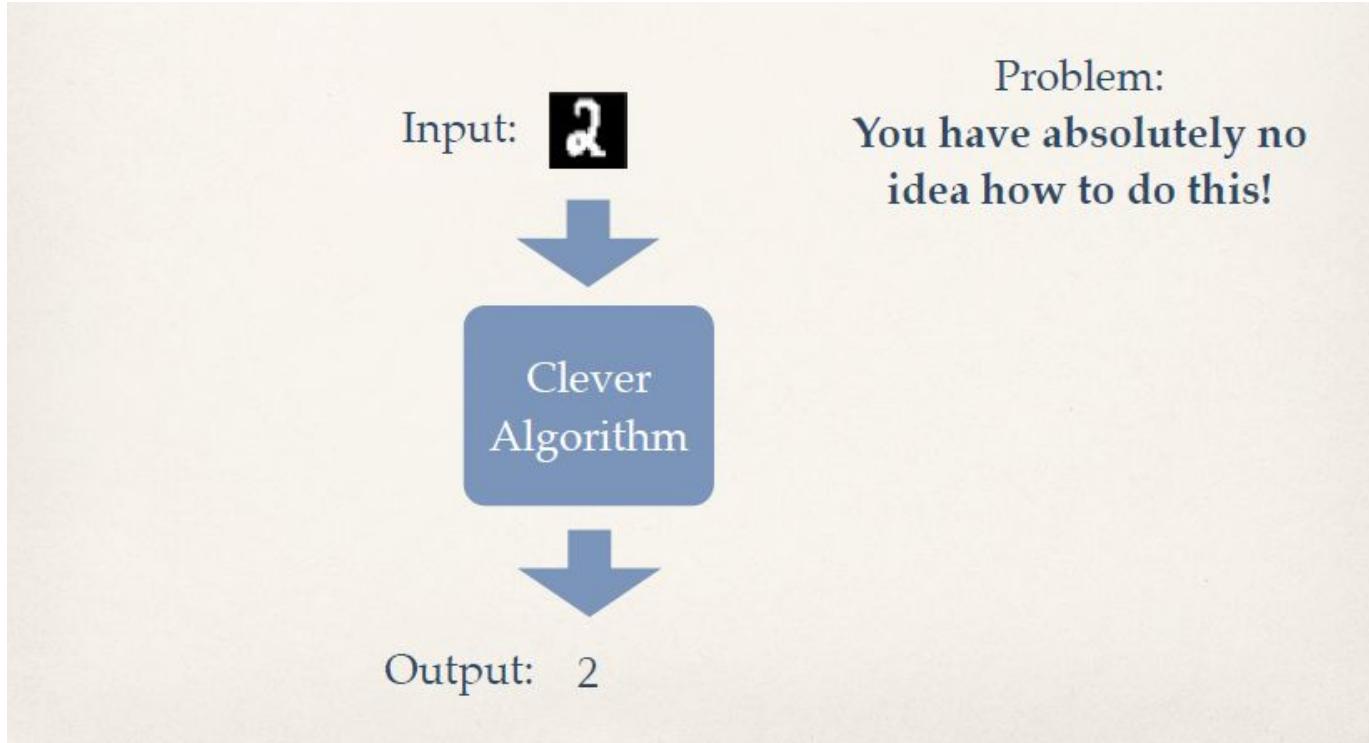
# Example: Spam email filtering

“A computer program is said to *learn* from experience E with respect to some task T and some performance measure P, if its performance on T, as measured by P, improves with experience E.”

- Suppose your email program watches which emails you do or do not mark as spam, and based on that learns how to better filter spam.
- Task:
  - Classifying emails as spam or not spam.
- Experience:
  - Watching you label emails as spam or not spam.
- Performance measure
  - The number (or fraction) of emails (disjoint from the training emails) correctly classified as spam/not spam.

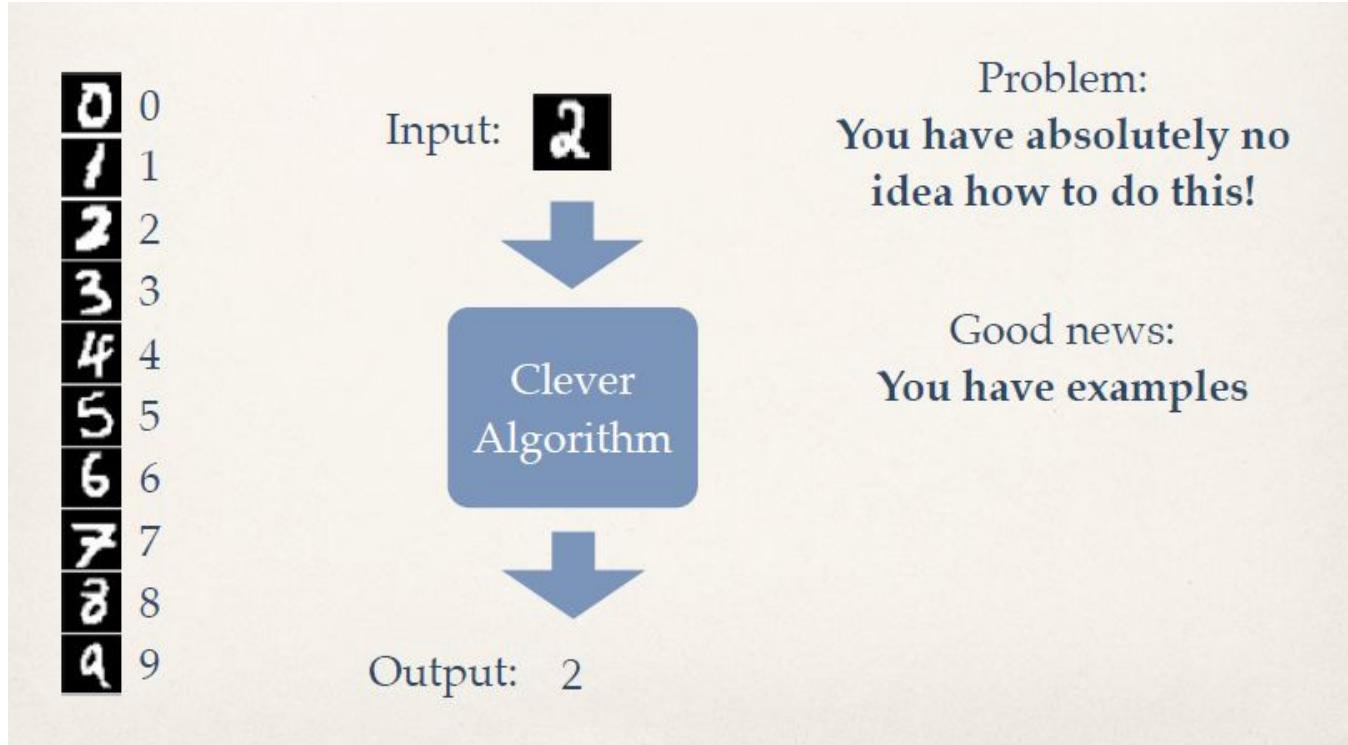
# Example

- Problem: Given an image of a handwritten digit, what digit is it?



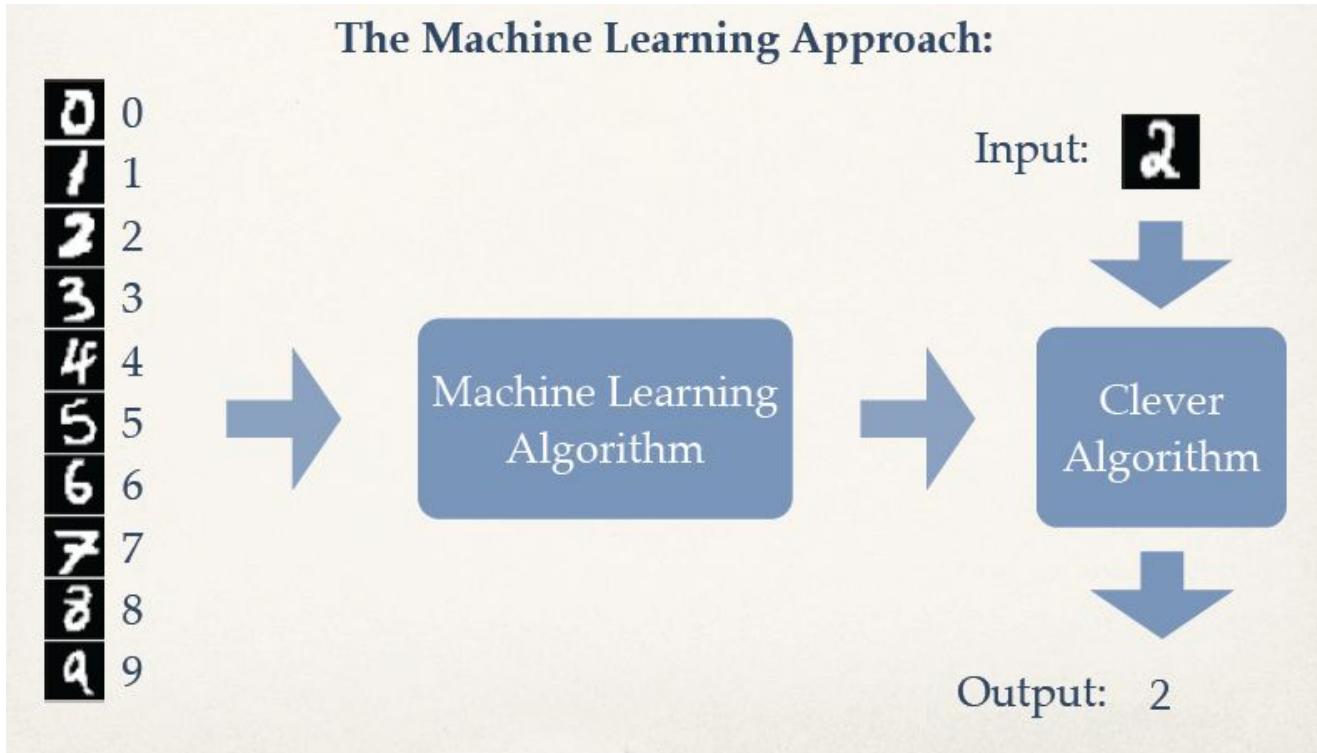
# Example

- Problem: Given an image of a handwritten digit, what digit is it?



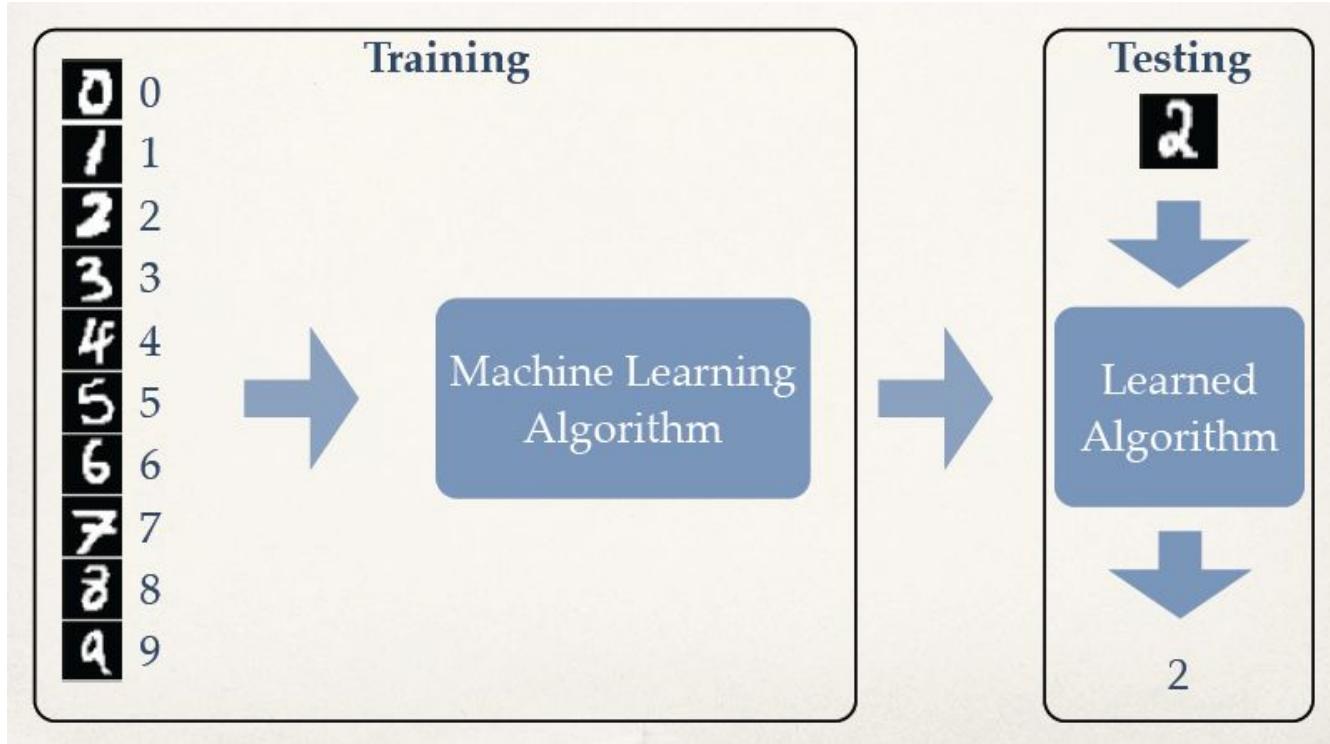
# Example

- Problem: Given an image of a handwritten digit, what digit is it?



# Example

- Problem: Given an image of a handwritten digit, what digit is it?



# Machine Learning Tasks

- Supervised Learning
  - Classification
  - Regression
- Unsupervised Learning
  - Clustering
  - Density estimation
  - Embedding / Dimensionality reduction
- Reinforcement Learning
  - Learning to act (e.g., robot control, decision making, etc.)

# Supervised Learning

Given a dataset  $D = \{(x_1, y_1), \dots, (x_n, y_n)\}$ ,  
where

- $x_i \in \mathcal{X}$  : input (feature)
- $y_i \in \mathcal{Y}$  : output (label)

a black box ML algorithm produces a  
prediction function  $h : \mathcal{X} \rightarrow \mathcal{Y}$ , such that  
 $h(x)$  can predict the  $y$  values for all  $x$   
(including training data  $x_i \in D$  and unseen  
test data  $x^*$  ).

# Supervised Learning

- Labels could be discrete or continuous
  - Discrete labels: **classification**
  - Continuous labels: **regression**

# Supervised Learning - Classification

Feature Space  $\mathcal{X}$



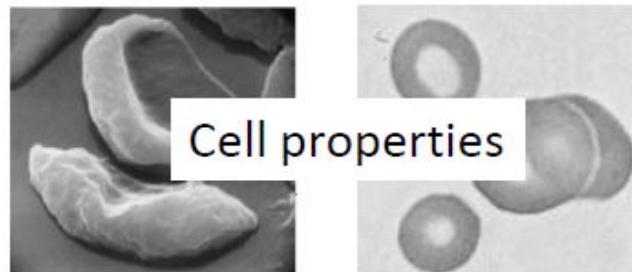
Label Space  $\mathcal{Y}$

- “Sports”
- “News”
- “Science”

...



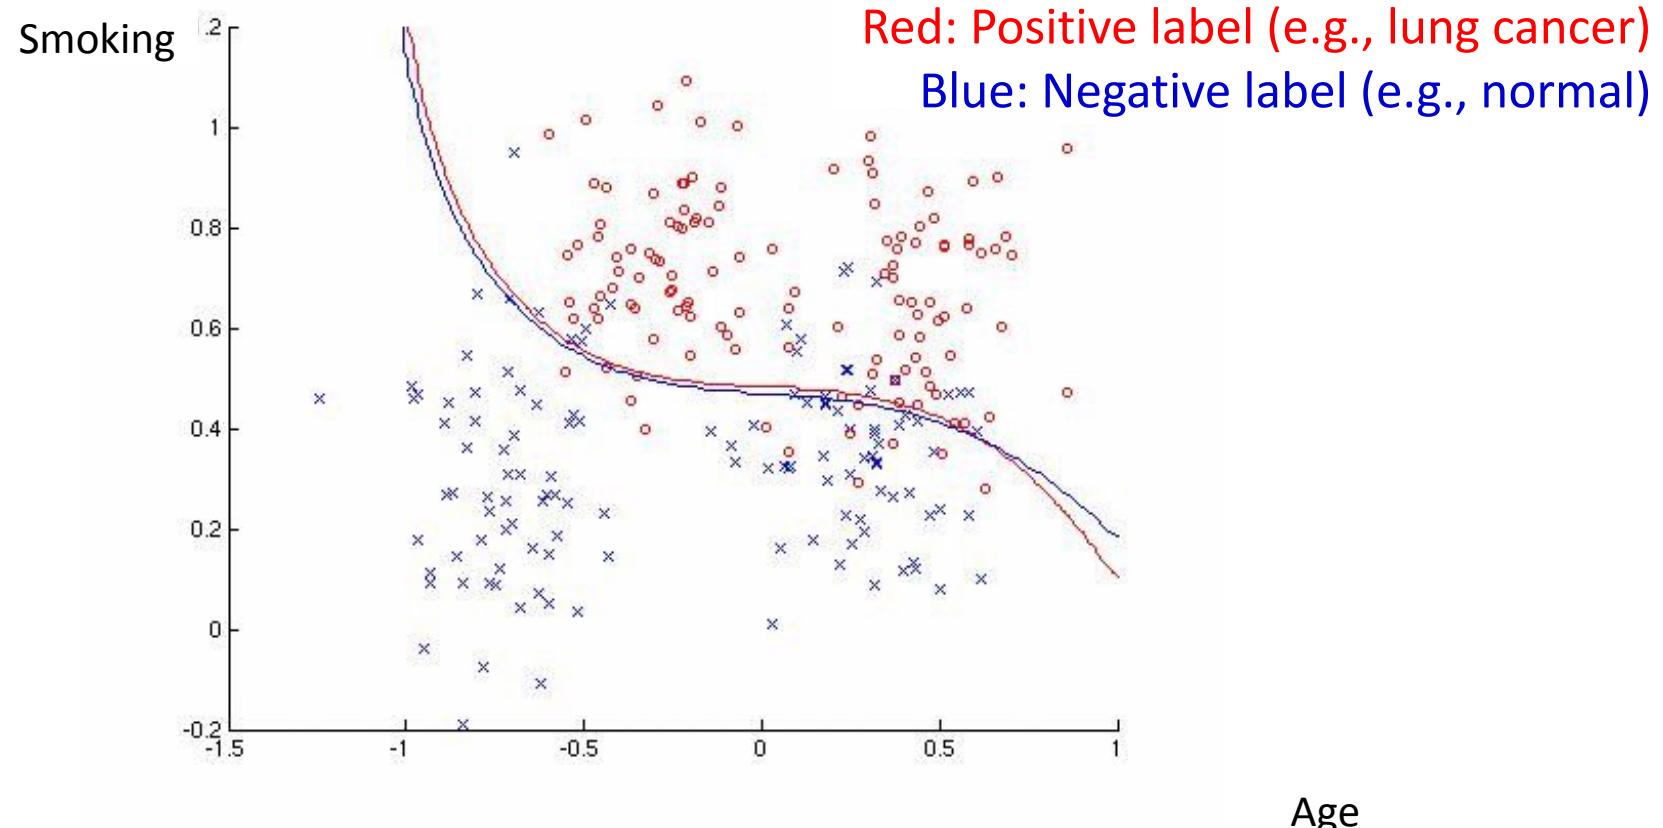
Cell properties



- “Anemic cell”
- “Healthy cell”

Discrete Labels

# Supervised Learning - Classification



“Learning decision boundaries”

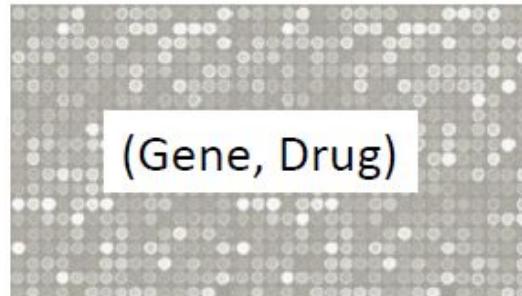
# Supervised Learning - Regression

Feature Space  $\mathcal{X}$



Label Space  $\mathcal{Y}$

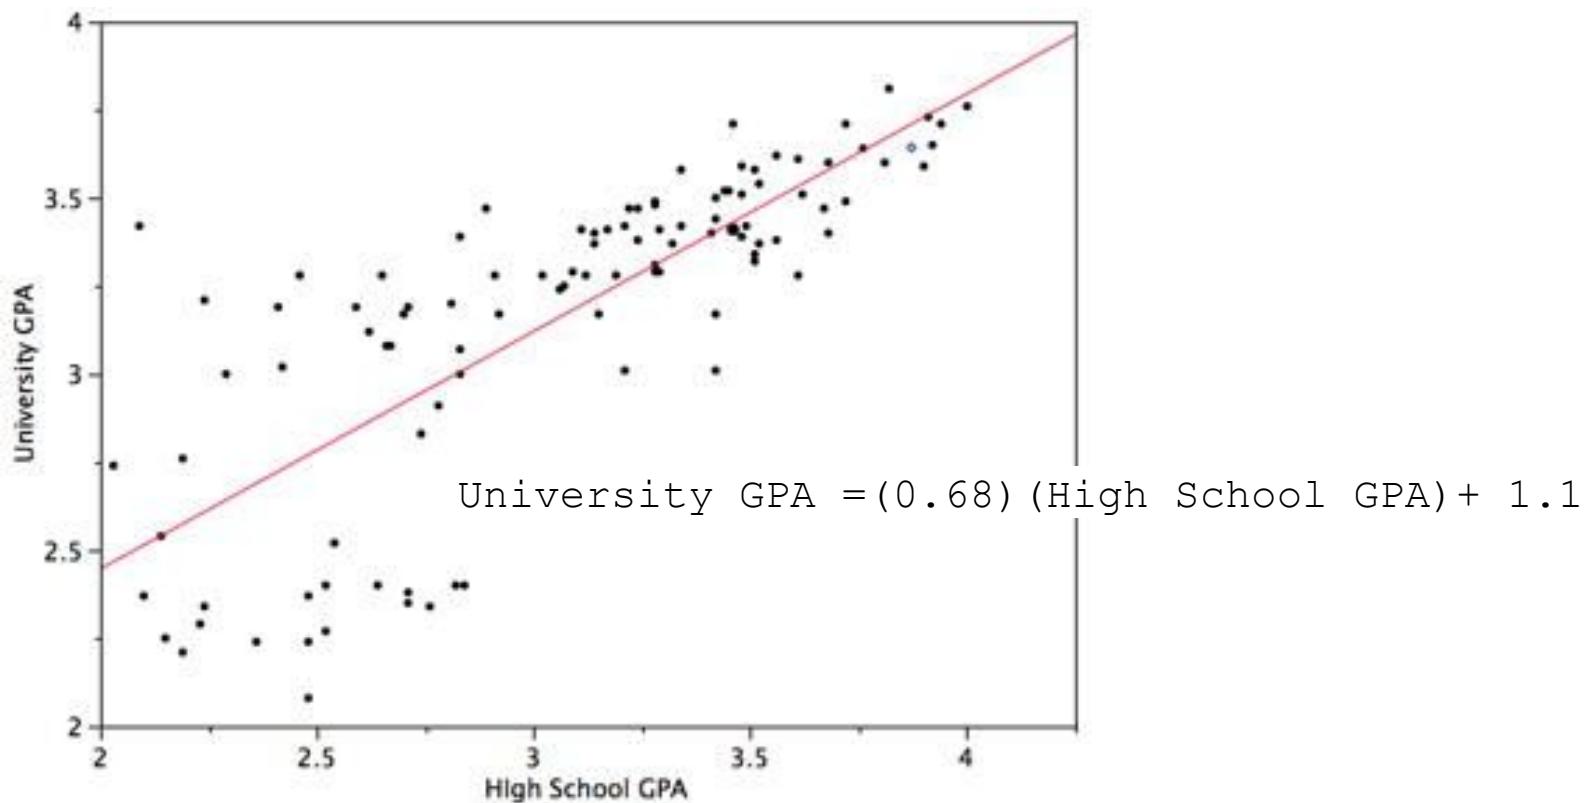
Share Price  
“\$ 24.50”



Expression level  
“0.01”

**Continuous Labels**

# Supervised Learning - Regression



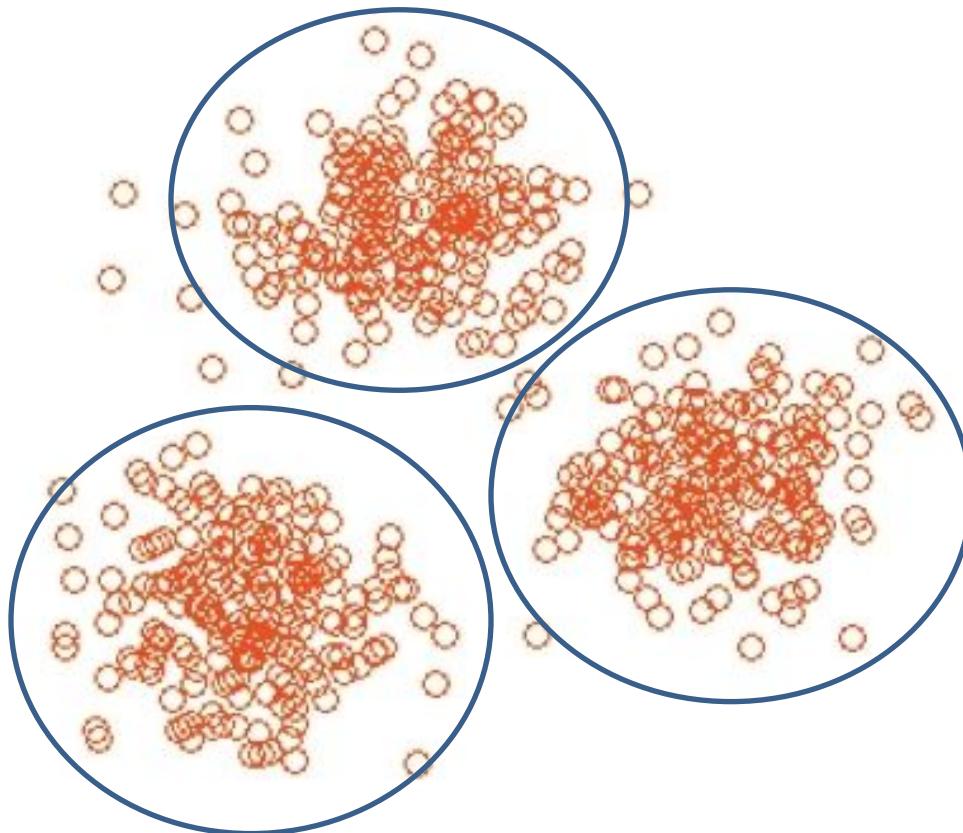
“Learning regression function  $f(X)$ ”

# Unsupervised Learning

- Goal:
  - Given data  $X$  without any labels
  - Learn the **structures** or **distribution** of the data
    - Clustering
    - Probability distribution (density)
    - Generating data
    - Embedding & neighborhood relations
- “Learning without teacher (supervision)”

# Unsupervised Learning – Clustering

- “Grouping into similar examples”



# Unsupervised Learning – Clustering

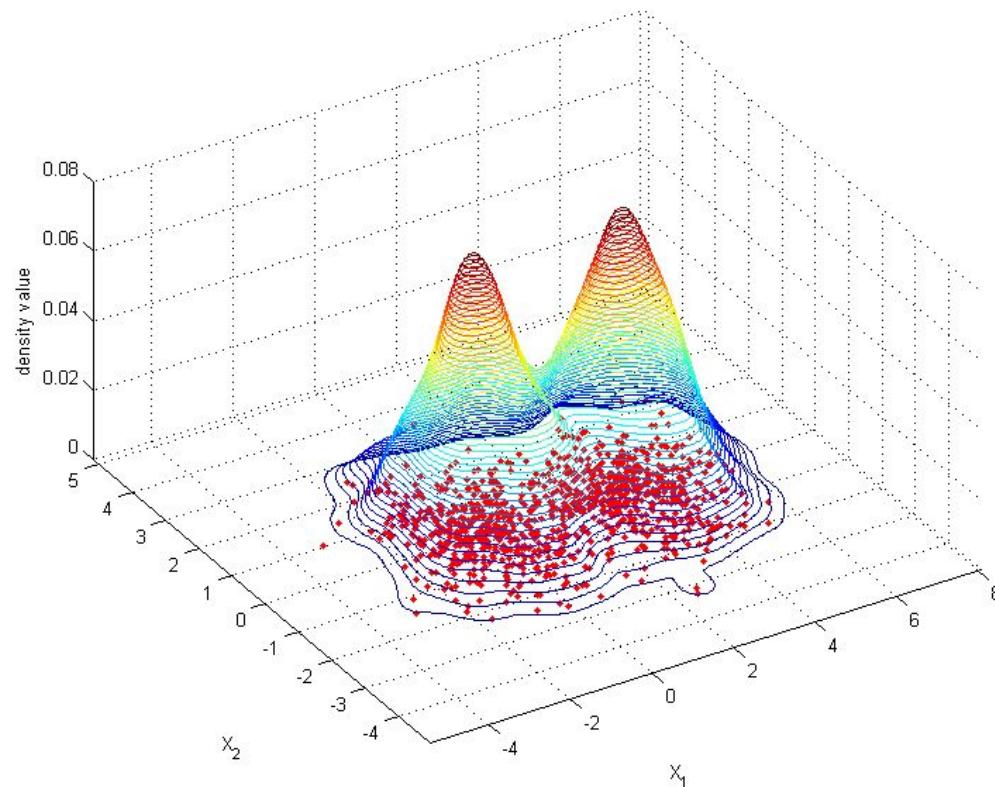
Group similar things e.g. images

[Goldberger et al.]



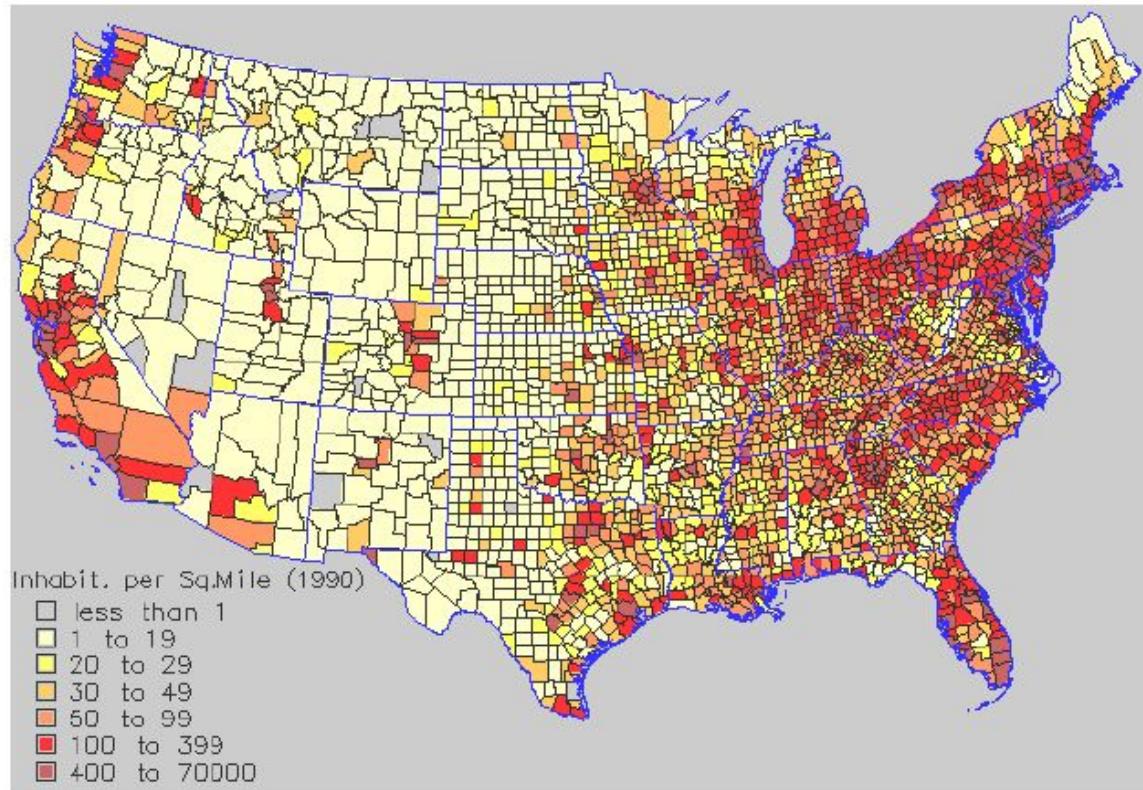
# Unsupervised Learning – Density estimation

$P(X_1, X_2)$   
“Probability”



# Unsupervised Learning – Density estimation

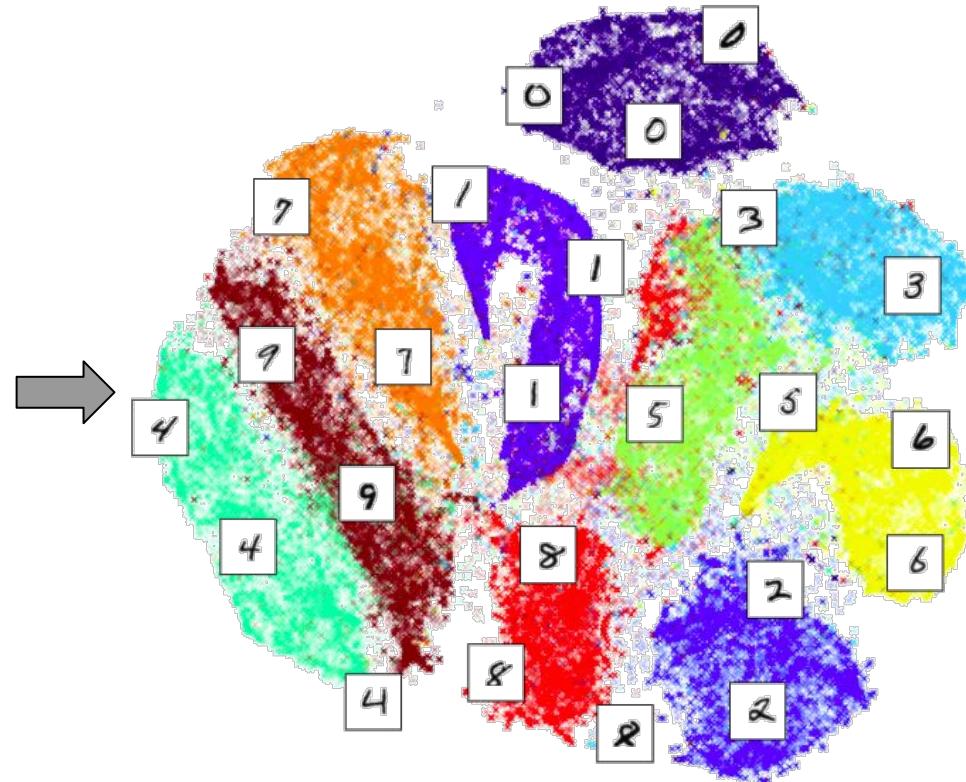
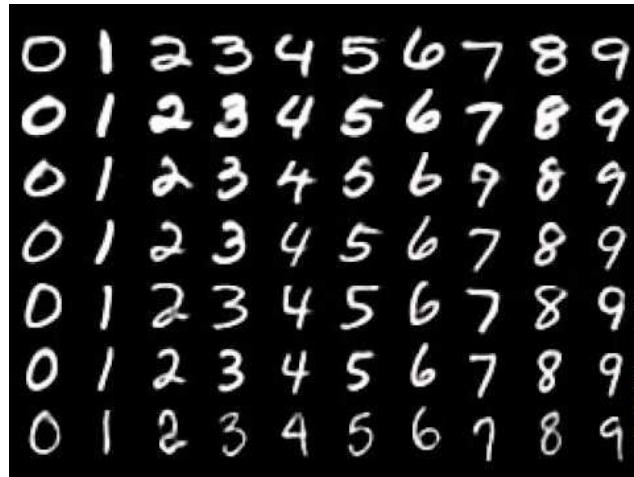
Population density



Slide credit: Aarti Singh

# Unsupervised Learning-Embedding and Dimensionality reduction

- E.g., Reducing handwritten digits (784 dim) into low dimensional coordinates



[Maaten and Hinton, 08]

# Reinforcement Learning

- Setting
  - Given sequence of states X and “rewards” (e.g., delayed labels)
  - Agent has to take actions A for each time step
- Goal:
  - How to “learn to act” or “make decisions” to maximize the sum of future rewards
- Example: Robot navigation task
  - Input: Dynamical environment + sensor input
  - Action: control signals
  - Rewards: time to reach goal without colliding with obstacles

# Reinforcement Learning – learning to control

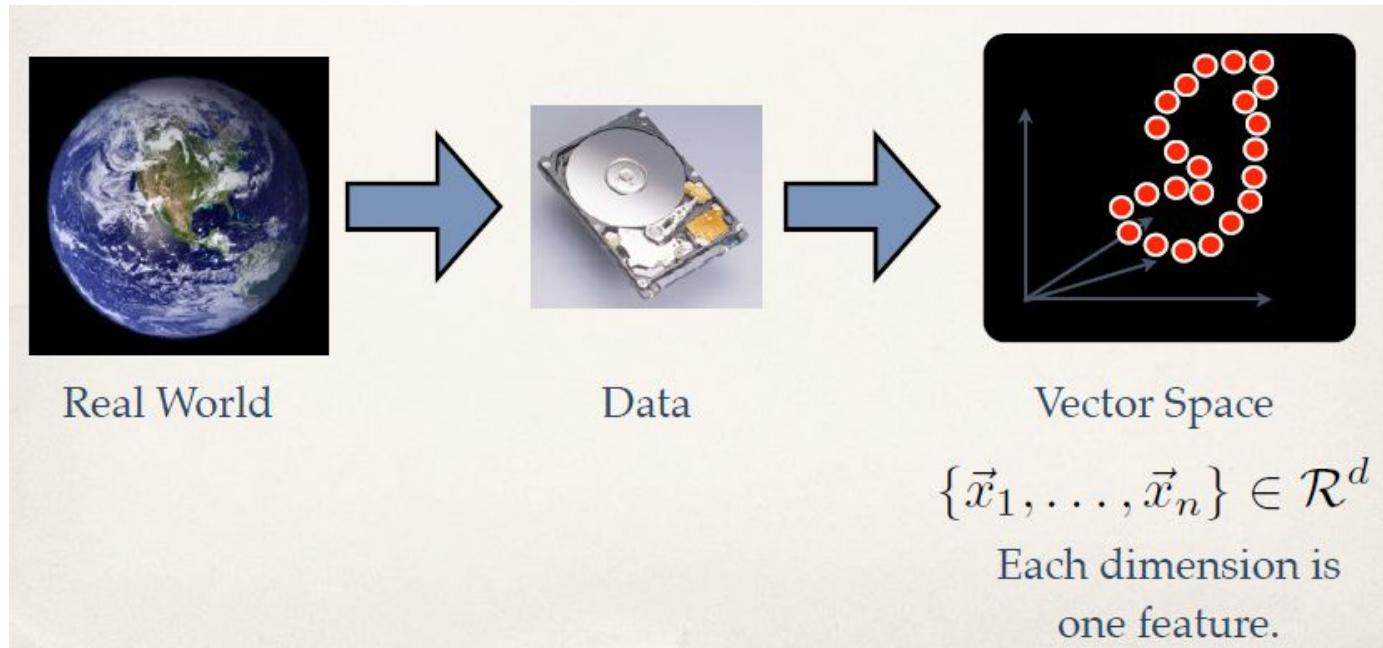
- Example: Robot walking
  - States: sensor inputs, joint angles
  - Action: servo commands for joints
  - Rewards:
    - 1 for reaching the goal
    - -1 for falling down
    - 0 otherwise
- Goal: How can we provide control inputs to maximize the expected future rewards?



# Feature representations

# Feature Extraction

- Represent data in terms of vectors.
  - Features are **statistics** or **attributes** that describe the data.



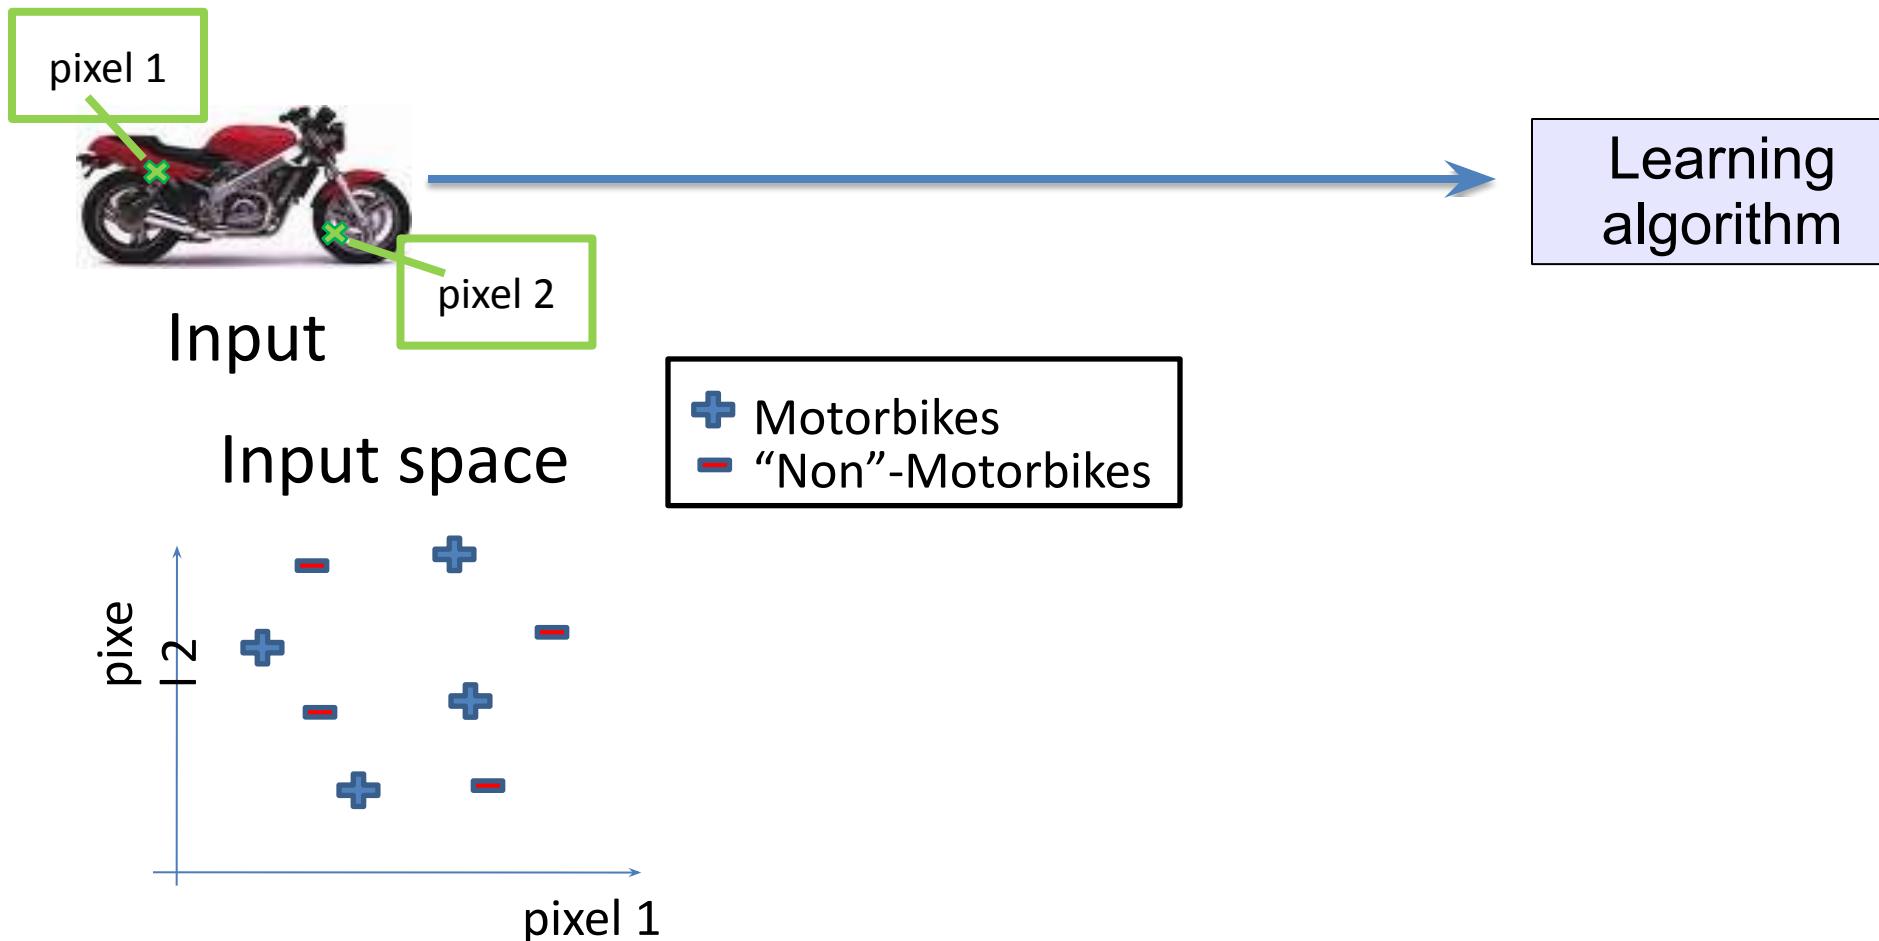
# Examples of features: Housing data

- Given statistics about houses in a local area, predict median value of homes.
  - #ROOM: average number of rooms per dwelling
  - AREA: average area of house in square foot
  - AGE: proportion of owner-occupied units built prior to 1940
  - CRIME: per capita crime rate by town
  - RESZONE: proportion of residential land zoned for lots over 25,000 sq.ft.
  - INDUS: proportion of non-retail business acres per town
  - NOX: nitric oxides concentration (parts per 10 million)
  - .....
- Label: Median value of owner-occupied homes

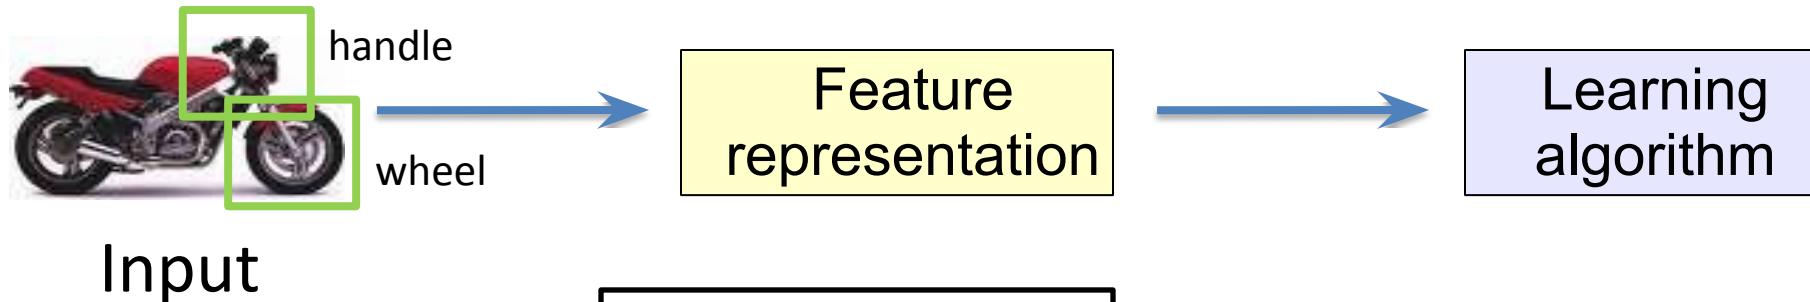
# Examples of features: Recognizing handwritten-digits

- Input:  $28 \times 28$  pixel images
- Output: Class Labels  $\in \{0, 1, 2, \dots, 9\}$
- The following basic features can be used:
  - Pixel Values (784 dimensional vectors)
  - Aspect Ratio of the tight bounding boxes
  - Existence of long vertical strokes
  - Existence of long horizontal strokes

# Learning pipeline

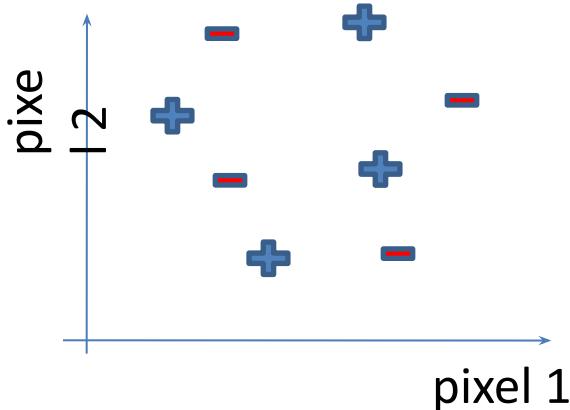


# Learning pipeline

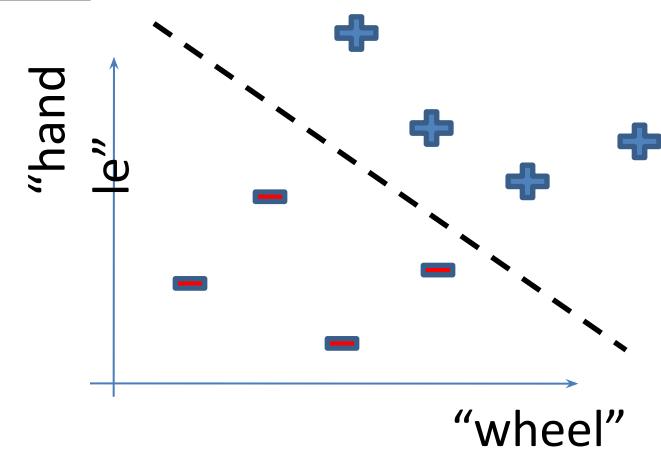


Input

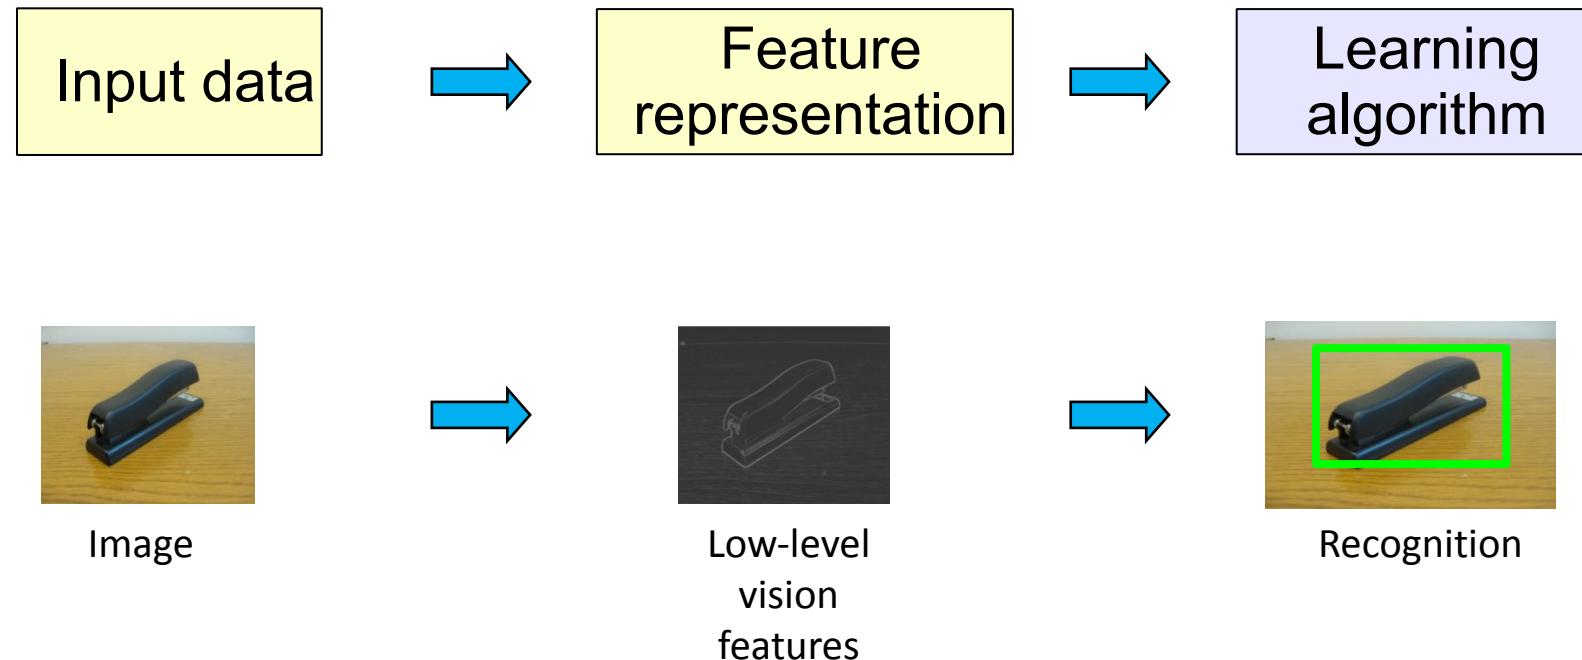
Input space



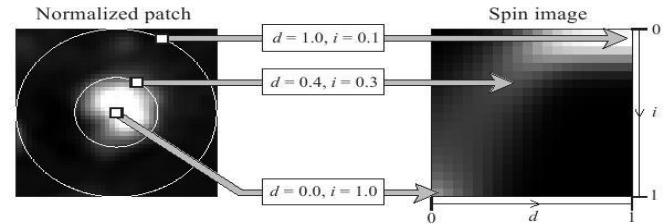
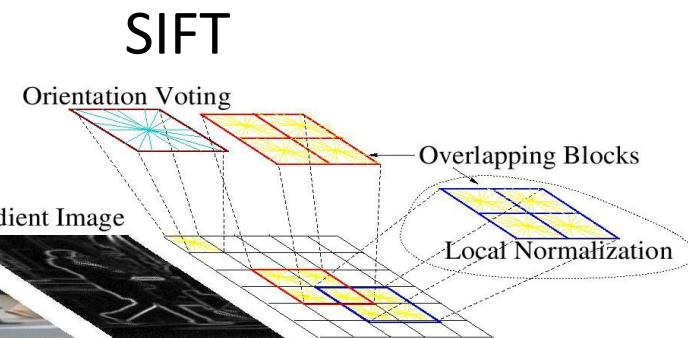
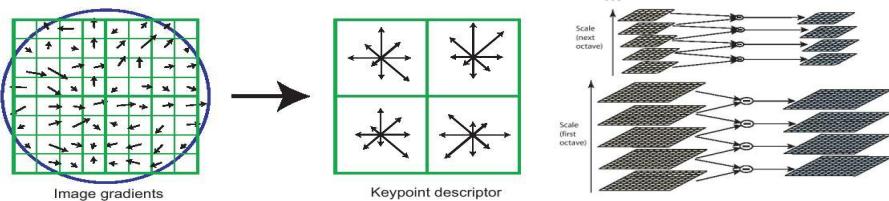
Feature space



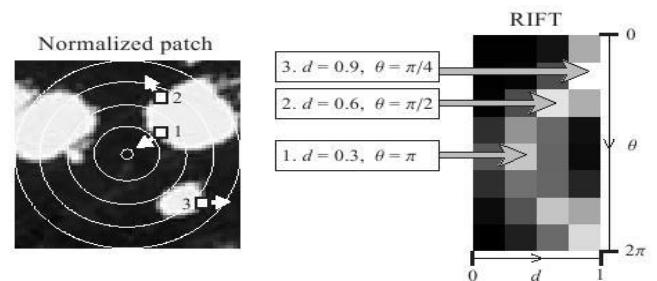
# (Traditional) Computer Perception Pipeline



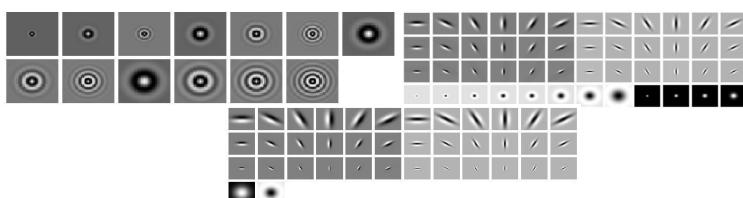
# (Traditional) Computer vision features



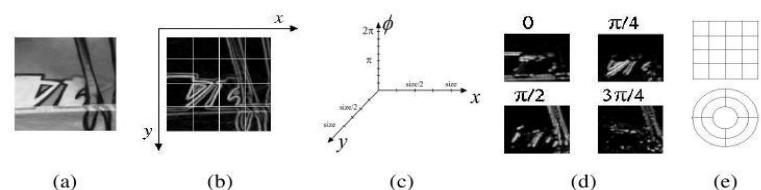
### Spin image



### RIFT



### Textons



### GLOH

# Learning feature hierarchies with Deep Learning

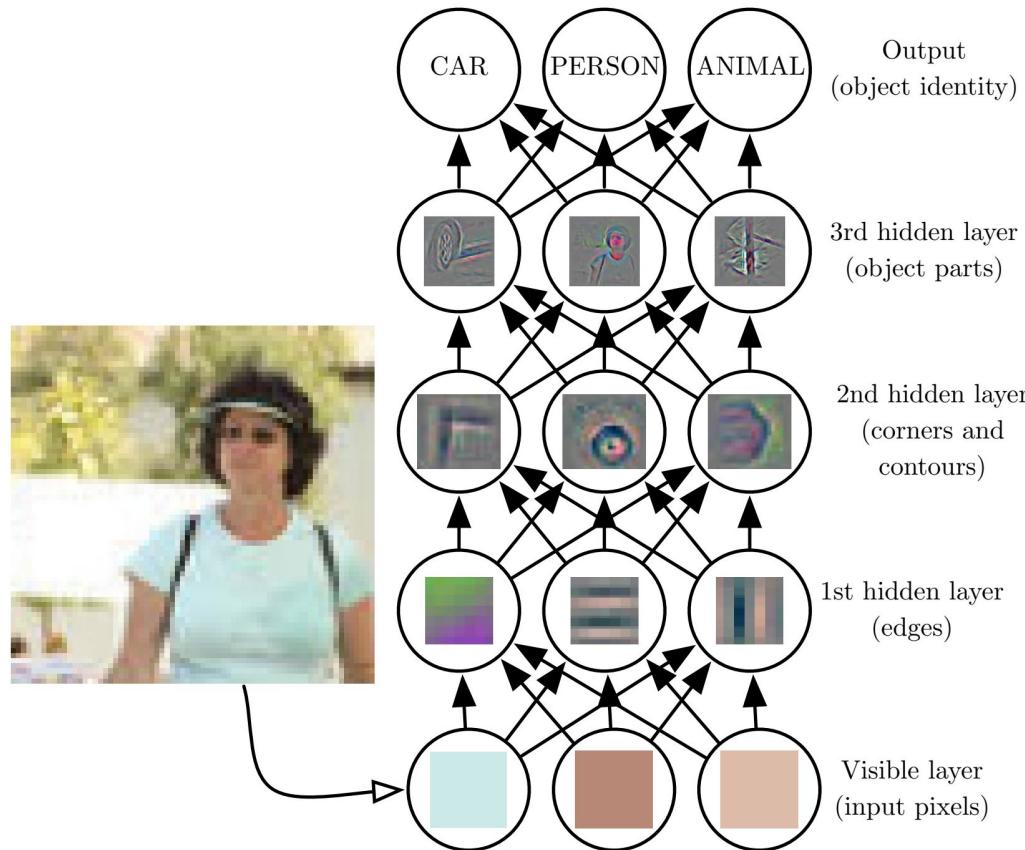
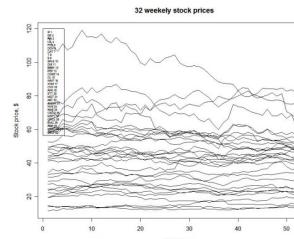
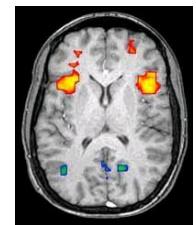
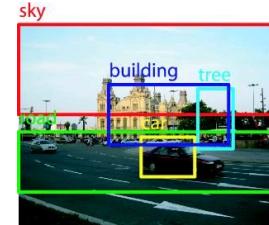


Figure source: Ian Goodfellow, Yoshua Bengio, Aaron Courville. Deep Learning, 2016

# ML applications

# Examples of ML applications

- Computer vision
- Speech recognition
- Robotics
- Text classification
- Medical image recognition
- Time series prediction/classification
- ....



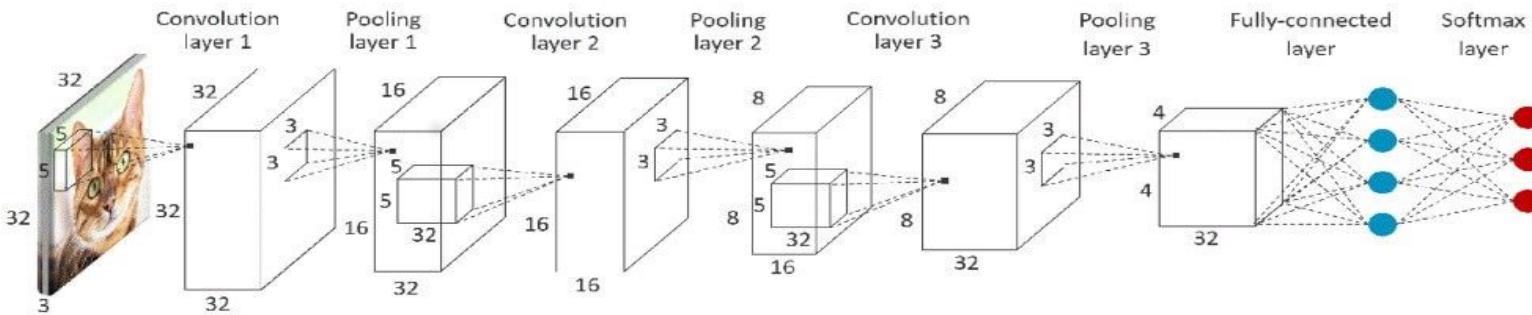
# Image classification

- Goal: predict the class of input image

Training dataset  
(image, label)



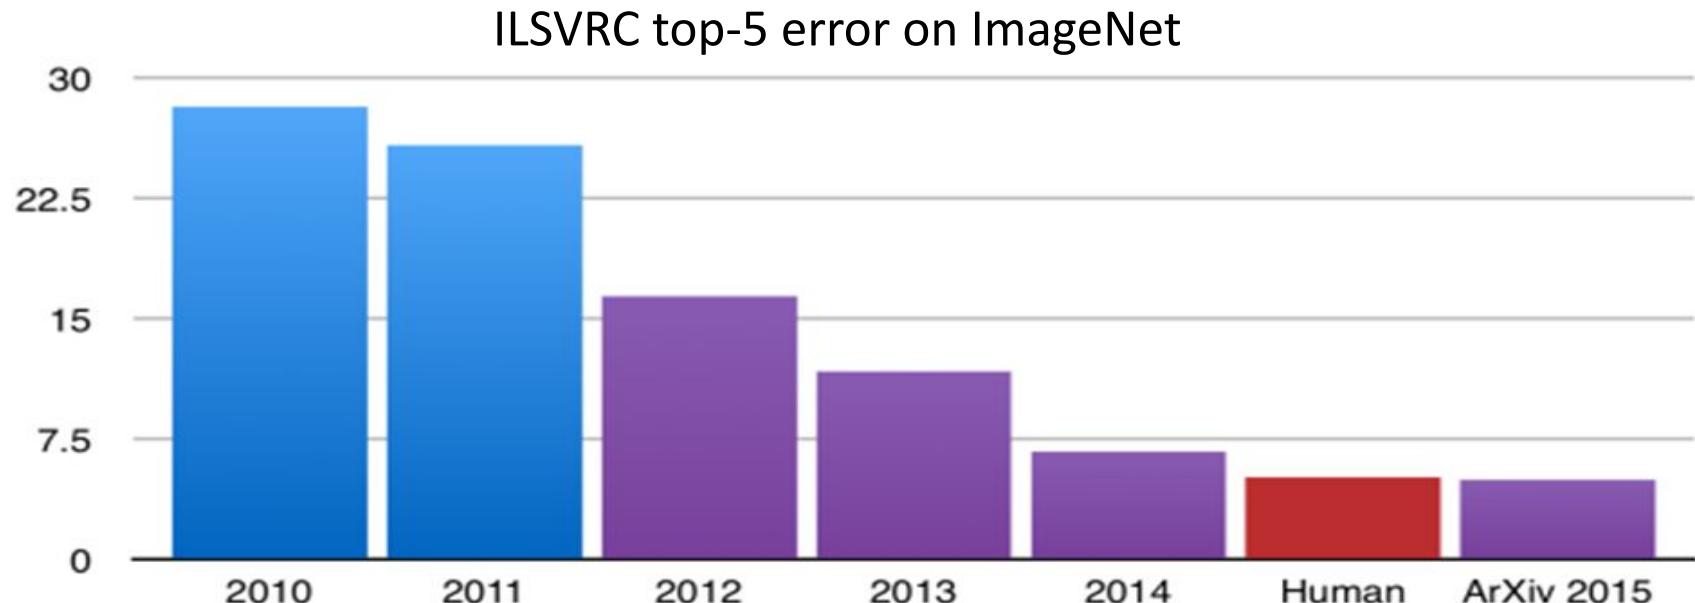
Input



Prediction

cat

# Image classification



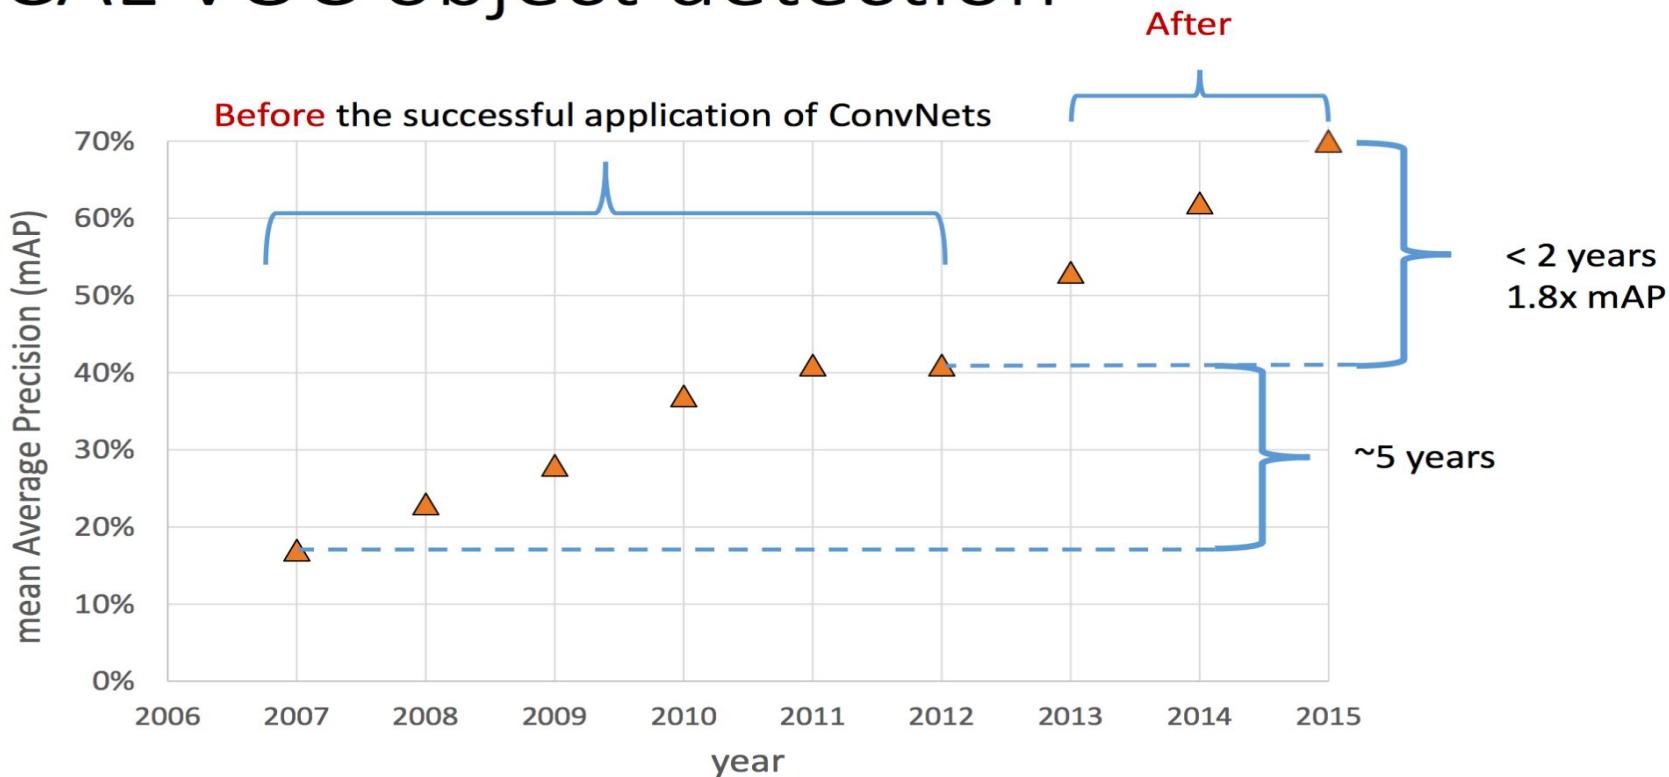
Slide credit: Rob Fergus

# Object Detection



Huang et al., Speed/accuracy trade-offs for modern convolutional object detectors. CVPR 2017  
figure/code at: [https://github.com/tensorflow/models/tree/master/research/object\\_detection](https://github.com/tensorflow/models/tree/master/research/object_detection)

# PASCAL VOC object detection

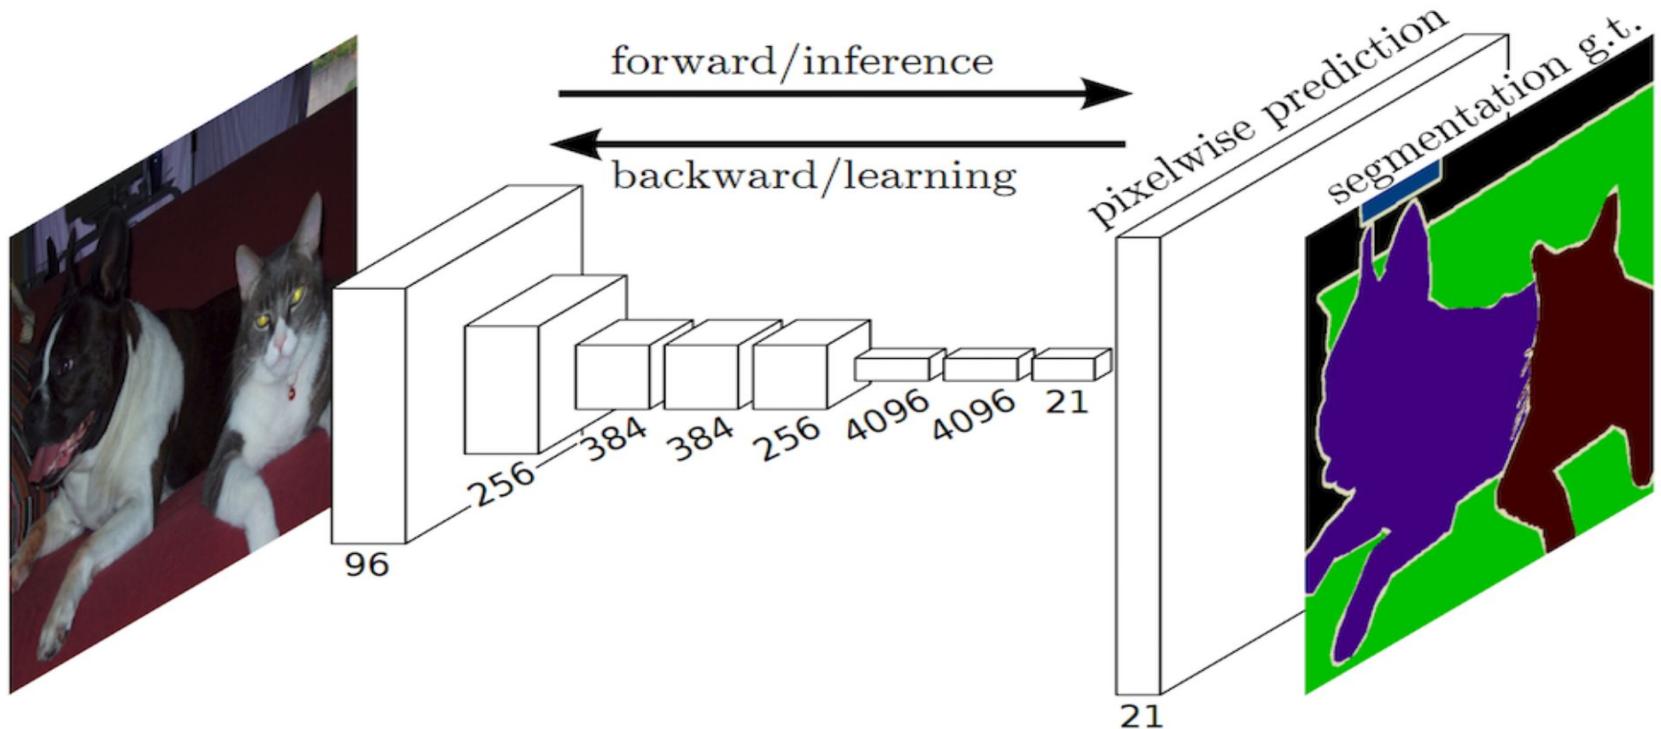


Precision: higher is better

(Figure from Ross Girshick)

# Semantic Segmentation

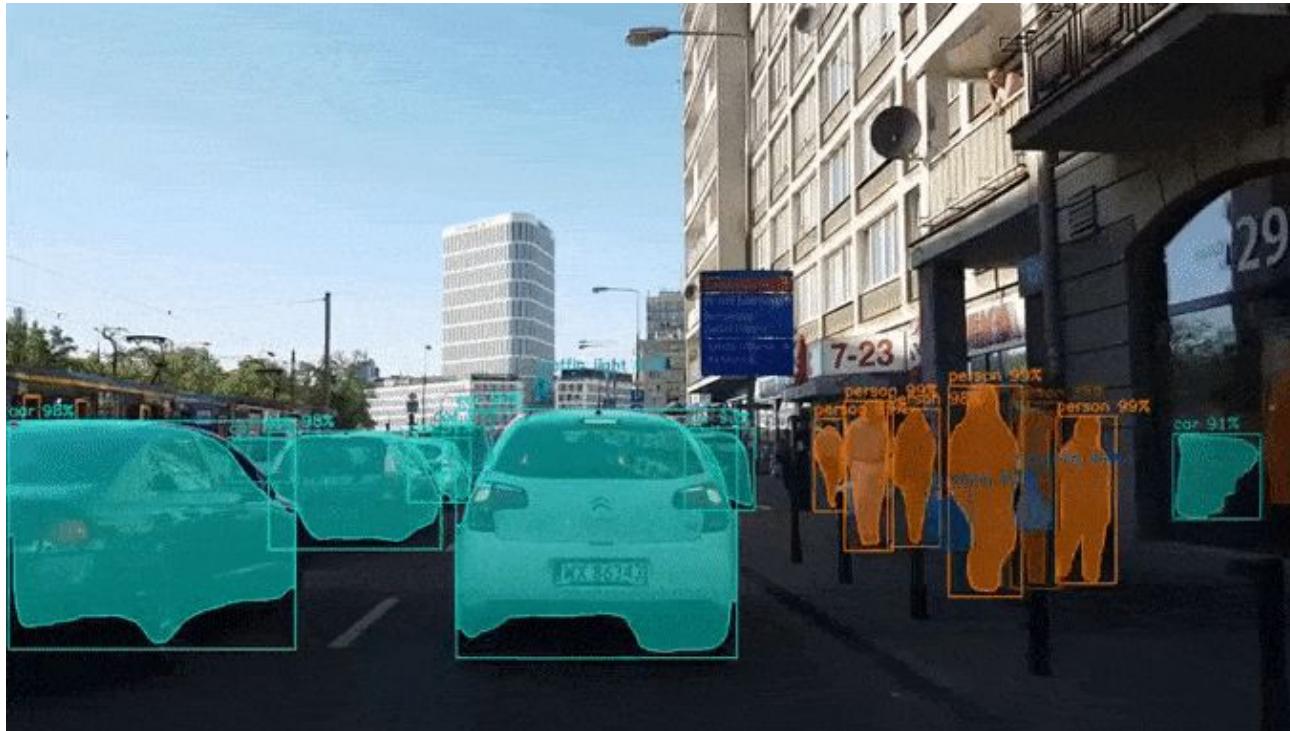
- Goal: Segment object regions and predict class labels for each region
- Can be formulated as pixel-wise classification



(Long et al, "Fully Convolutional Networks for Semantic Segmentation", CVPR, 2015.)

# Instance Segmentation

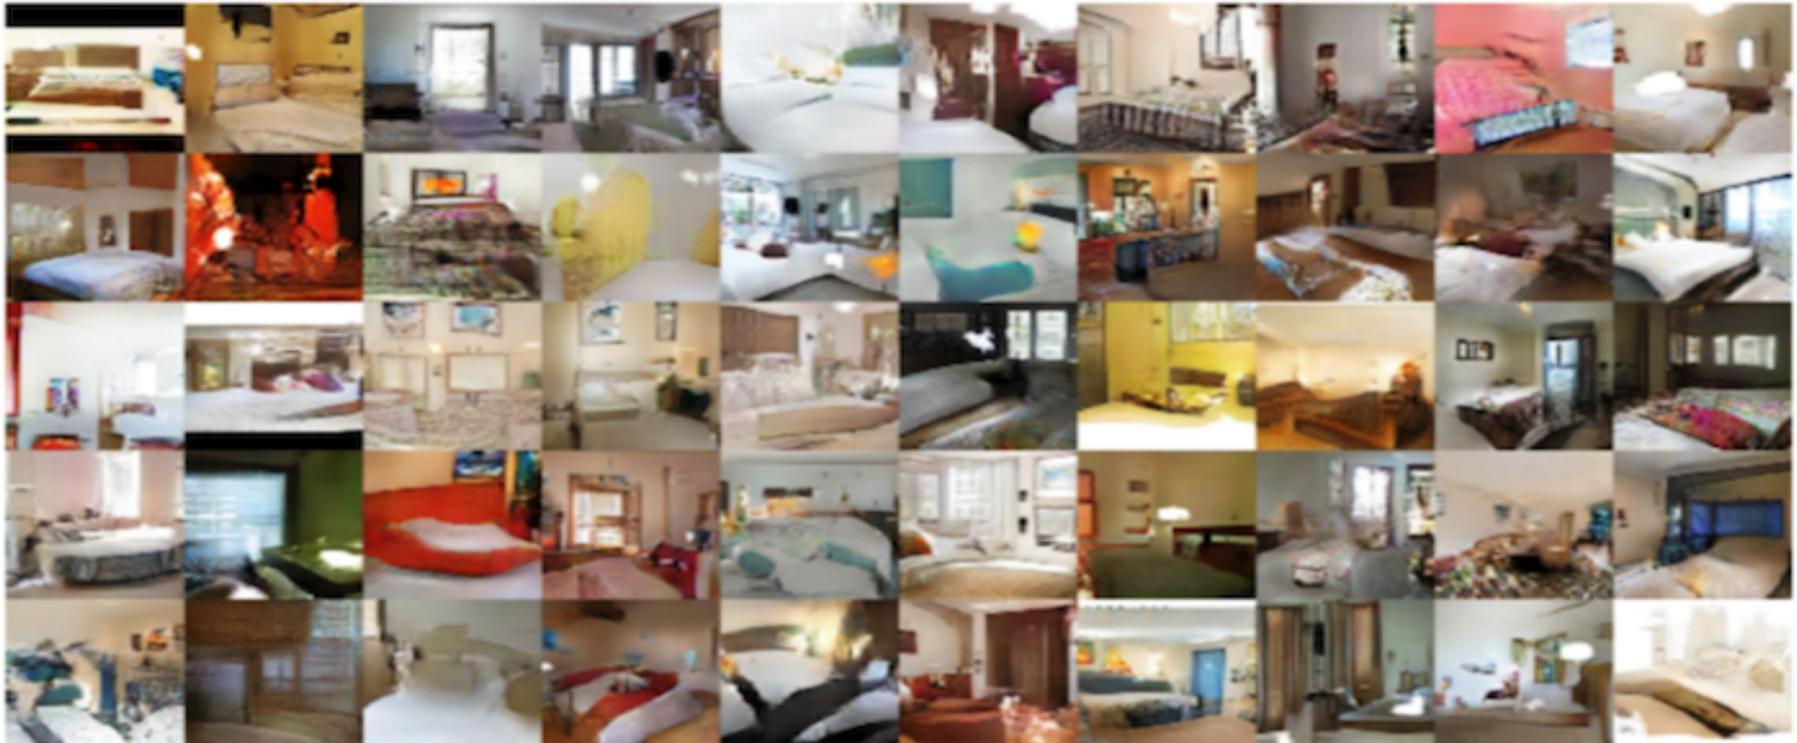
Goal: Perform segmentation (pixel-level masking) for each individual object



Mask R-CNN (He et al., 2017)

Image from: [https://github.com/matterport/Mask\\_RCNN](https://github.com/matterport/Mask_RCNN)

# Image Generation: Generative Adversarial Networks



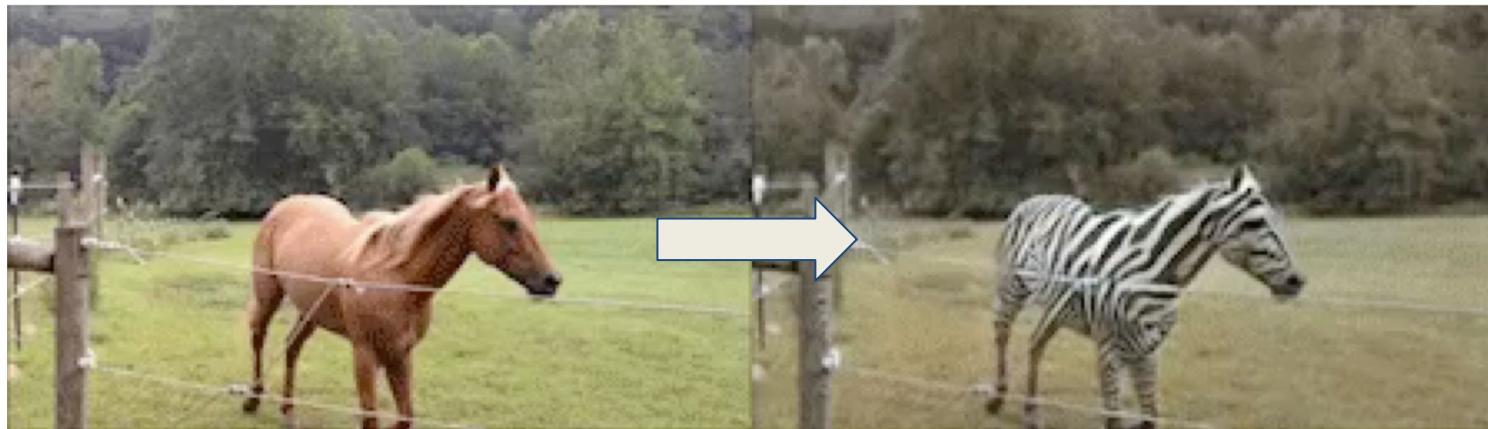
(Radford et al, "Unsupervised Representation Learning with Deep Convolutional Generative Adversarial Networks", ICLR, 2016.)

# Image Generation: Generative Adversarial Networks



[Karras et al. 2017] Progressive Growing of GANs for Improved Quality, Stability, and Variation

# Unsupervised Image-to-Image generation



# Image Caption Generation

A group of young people playing a game of frisbee.



A herd of elephants walking across a dry grass field.



Two hockey players are fighting over the puck.



A close up of a cat laying on a couch.



A little girl in a pink hat is blowing bubbles.



A red motorcycle parked on the side of the road.



A refrigerator filled with lots of food and drinks.



A yellow school bus parked in a parking lot.



Describes without errors

Describes with minor errors

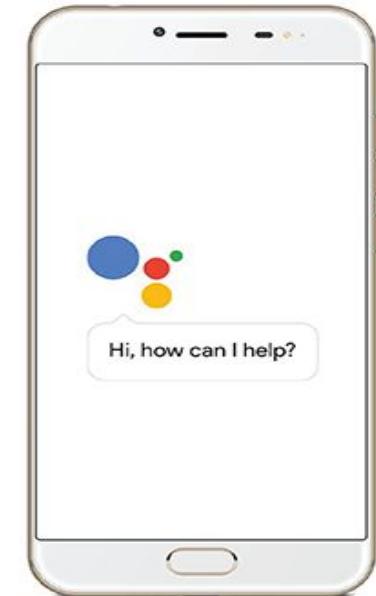
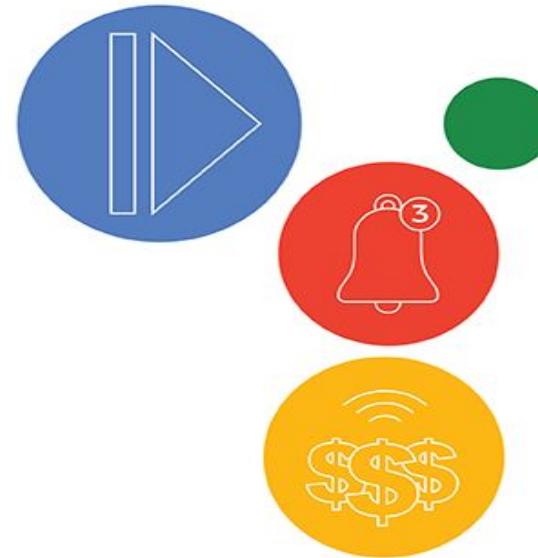
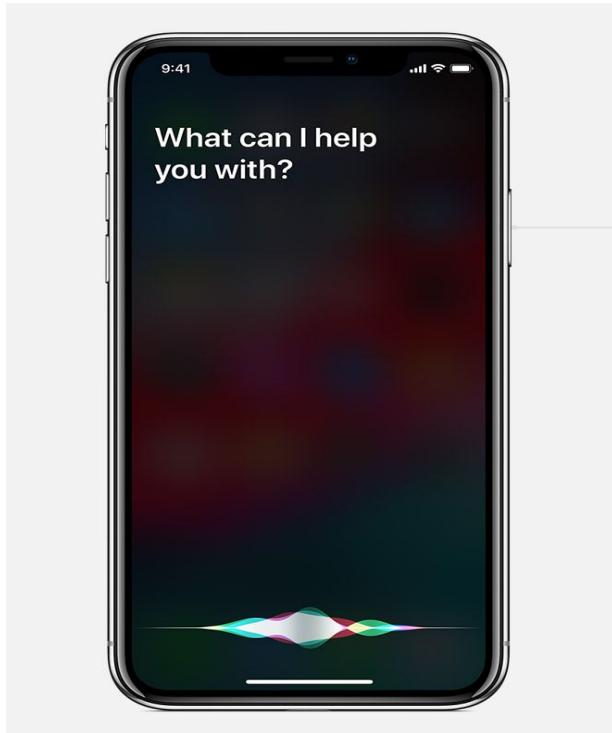
Somewhat related to the image

Unrelated to the image

(Vinyals et al, "Show and Tell: A Neural Image Caption Generator", CVPR, 2015.)

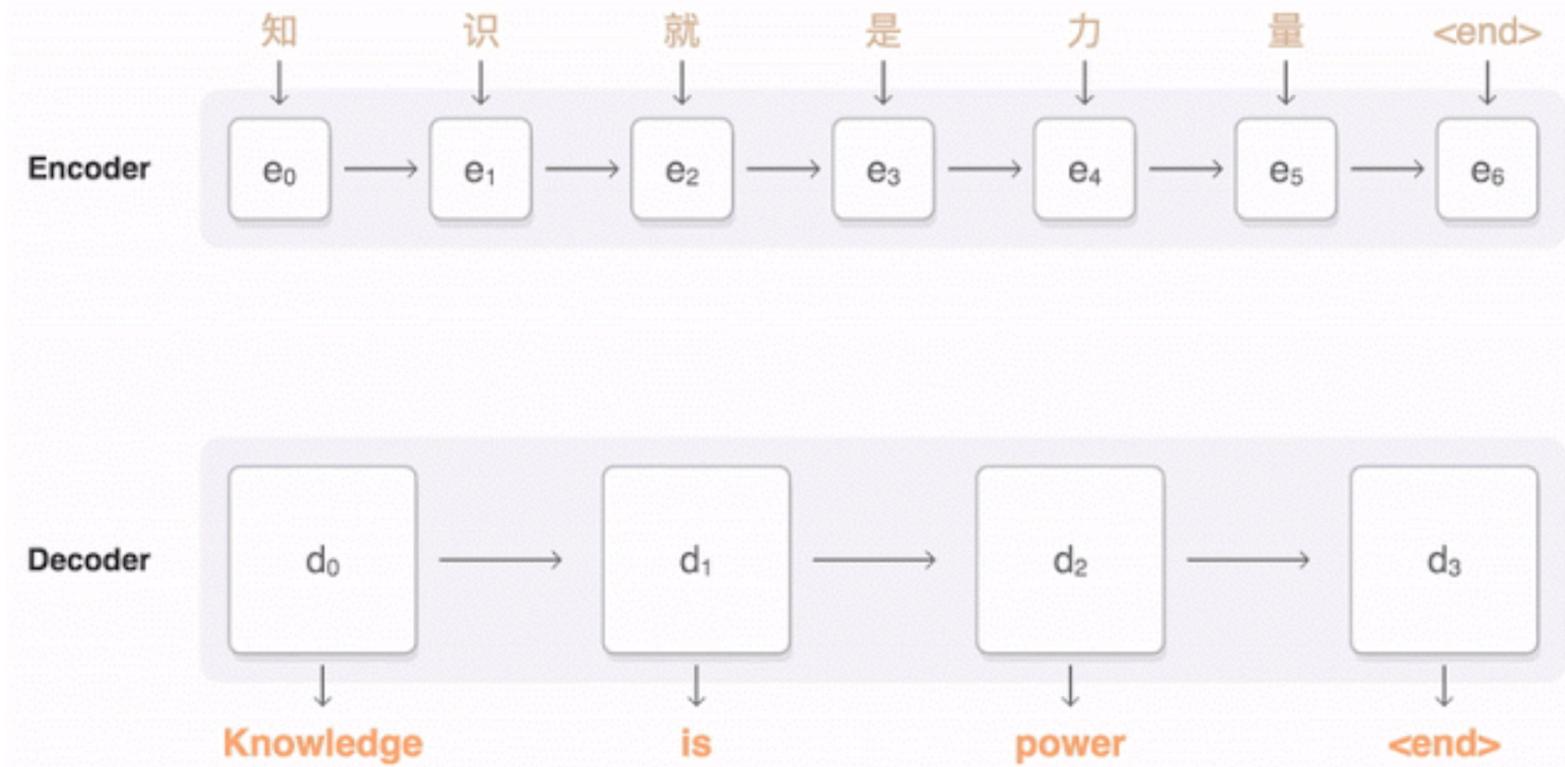
# Speech recognition

Siri, Google home, and Google assistant achieves commercial-level performance

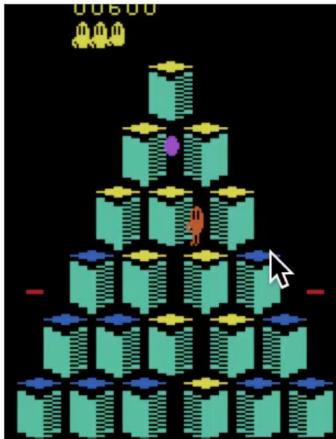


# Machine Translation

Google Neural Machine Translation (in production)



# RL success stories: playing ATARI games



**DQN** Mnih et al, NIPS 2013 / Nature 2015;

**MCTS** Guo et al, NIPS 2014; **TRPO** Schulman, Levine, Moritz, Jordan, Abbeel, ICML 2015;  
**A3C** Mnih et al, ICML 2016; **Dueling DQN** Wang et al ICML 2016; **Double DQN** van Hasselt et al, AAAI 2016; **Prioritized Experience Replay** Schaul et al, ICLR 2016; **Bootstrapped DQN** Osband et al, 2016; **Q-Ensembles** Chen et al, 2017; **Rainbow** Hessel et al, 2017; ...

# AlphaGo

- Another breakthrough from Google DeepMind
- Combines Monte-Carlo Tree Search (MCTS) with deep neural networks



**AlphaGo** Silver et al, Nature 2015

**AlphaGoZero** Silver et al, Nature 2017

**AlphaZero** Silver et al, 2017

Tian et al, 2016; Maddison et al, 2014; Clark et al, 2015

# OpenAI's 1v1 Dota [2017] and 5v5 [2018]

Super-human agent on a competitive game, enabled by

- Reinforcement learning
- Self-play
- Enough computation

Cooperation emerges



# Robot learning



Levine et al., Learning Hand-Eye Coordination for Robotic Grasping. 2016  
video: [https://www.youtube.com/watch?v=cXaic\\_k80uM](https://www.youtube.com/watch?v=cXaic_k80uM)

slide credit: Pieter Abbeel 71

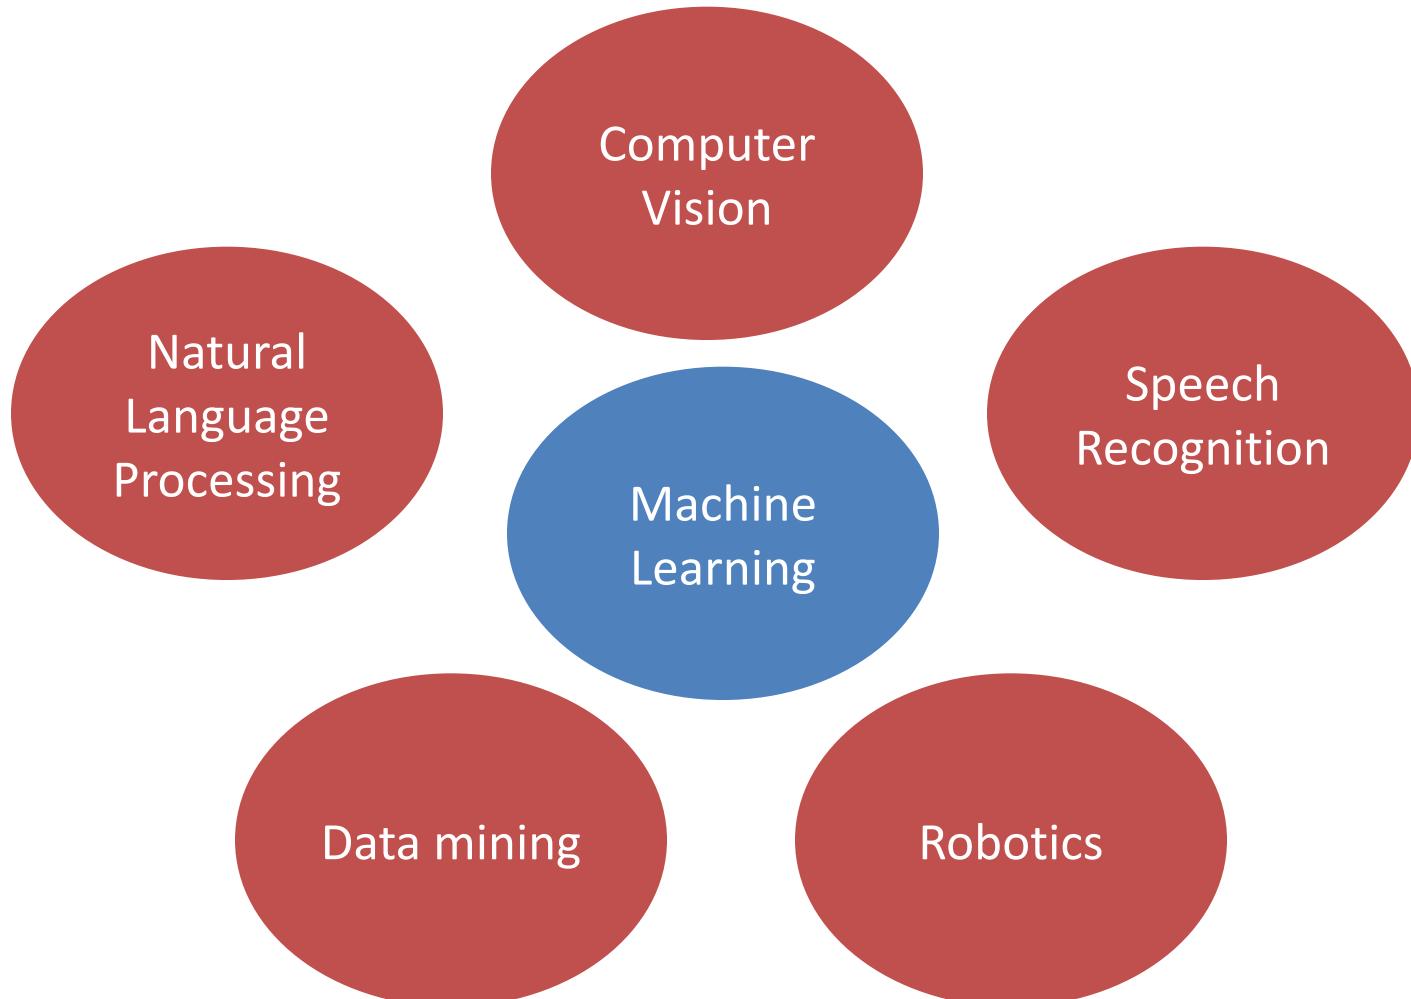
# Self-driving cars



<https://youtu.be/O6DRfAC1JXA>

See also: Chris Urmson: How a driverless car sees the road <https://youtu.be/tiwVMrTLUWg>

# Machine Learning and other fields



# Next class

- Supervised learning
  - Linear regression

# Reminder

- Check syllabus at Canvas
- For all questions, please use Piazza (linked to Canvas)

# Questions?